

CX TG Requirements Discussion #1: CX State Model: One State or Multiple States?

Jan Gray, Acting Vice-Chair, Composable Custom Extensions TG

RVIA Disclosures

- <https://wiki.riscv.org/display/HOME/Meeting+Disclosures>


Our first requirements discussion breakout

- 8:30: \leq 5 min position statements, hold questions pls
- Open discussion
- 9:25: summing up
- Sharing our understanding, towards a consensus decision
- How to decide?
 - Thoroughly understand design, utility, and cost up and down the stack
 - Weigh additional spec, POC, and ratification time vs. decades of use
 - Weigh cost to change course later
- No decisions today

Jan's position

- The CX software programming model's CX state model must support
 - Composition of independent CX software libraries
 - Modern programming models, incl. tasks, coroutines, async methods
 - Efficiently and uniformly
- The CX state model should be flexible for diverse use cases
 - 0, 1, few, #harts, many \times 0, S, M, L, XL CX state contexts
 - Isolated *and* shared state CX ISA design patterns
 - Bare metal and server class workloads and systems
- The basis spec's ' m harts : n CX state contexts' (CX instances as hardware objects) is more novel, more complex, but more flexible, more performant, affordable, preferable

Let us be crisp about abstraction layers



- **User programming model**
- User ISA
- Supervisor OS concerns
- Supervisor ISA
- Machine ISA
- Logic interface

Avoid kicking the complexity can down the road

Focusing on CX Software Programming Model

Single or multiple = not that different

- Once on library init

```
cx = cx_open(MYCX_GUID, sharing); // discover if CX present, allocate a state context
```

- Per client call

```
if (cx < 0) return sw(...); // pure software if CX absent
```

```
cx_select(cx); // csrwr cx_index,cx // MYCX custom instructions
```

```
for (...)
```

```
    c[i] = cx_reg(FN,a[i],b[i]); // custom-0 rd,FN,rs1,rs2
```

```
cx_select(0); // csrwr cx_index,zero// built-in custom instructions
```

```
... = cx_status(); // csrrr ...,cx_status
```

Multiple instances of things?

- Programmers want multiple instances of resources
 - My library uses it, your library uses it too
 - My task/coroutine uses it, my other does too
- In software we create multiple ADTs or objects or resource handles
 - Singletons are awkward
- Same, if programming with accelerated software libraries?

Multiple instances of accelerators?

- In basis spec, built into the selector and CX mux'ing
 - `cx_select(cx);` // selects CX and CX's state context, one instruction
 - Optional
- In lieu of this: design CX state & behavior for shared single instance
 - No suspended stateful CX interactions: “get in, do it, and get out”
 - Repeatedly load CX context, compute, save CX context – not uniform
 - One CX state context but with ad hoc per-CX sub-context IDs
- Still need multiples, but we build them ad hoc

CX instances as hardware objects

- A uniform way to support multiple accelerator instances
- Key differences
 - Potential for CX software to open a second instance of a stateful CX
 - Programmer regards as an owned resource
 - Owned, not a constituent of, hart's architectural state
 - Like file descriptors, C++ objects, etc., not saved in `makecontext()`

Jan's position

- The CX software programming model's CX state model must support
 - Composition of independent CX software libraries
 - Modern programming models, incl. tasks, coroutines, async methods
 - Efficiently and uniformly
- The CX state model should be flexible for diverse use cases
 - 0, 1, few, #harts, many \times 0, S, M, L, XL CX state contexts
 - Isolated *and* shared state CX ISA design patterns
 - Bare metal and server class workloads and systems
- The basis spec's ' m harts : n CX state contexts' (CX instances as hardware objects) is more novel, more complex, more flexible, more performant, affordable, preferable

Other positions

BACKUP

What is this TG about?

- Composition is our paramount concern
 - Prize #1: a composable extensions ecosystem with a catalog of mix-and-match reusable components
 - CX software libraries that issue CX custom instructions to CXUs
 - “What works separately, works together”
- Uniformity is very important too
 - Prize #2: a good enough, common way to do common CX stuff
 - So you don’t have to reinvent everything
 - So my way is also your way, so our ways are compatible
 - Here, uniform state model, OS support, and programming model

Some other design tenets [basis]

- Diversity
 - HW: simple ... complex CPUs, CXs/CXUs, topologies
 - E.g., $k \text{ clusters} \times \{ m \text{ harts} \times n \text{ CXs} \times [s_0, \dots, s_{n-1}] \text{ CX state contexts} \}$
 - SW: bare metal ... RTOS ... OS ... HV+OSs+VMs/JITs+apps
 - One trusted org @ one time ... dozens of random orgs @ many years
- Longevity
 - (Here) support diverse programming models, legacy, modern, emerging
- Simple, frugal, fast

Stateless and stateful CX custom instructions [basis]

- Basis spec requires & achieves invariance under composition
- Stateless CX: “each [insn] is a pure function of its operands”
- Stateful CX: “each [insn] may access, and as a side effect, update, the hart’s current state context of the extension (only)”
 - “The behavior ... only depends upon the series of [insns] ...”
 - “Besides updating extension state, ..., and a destination register, an [insn] has no effect upon any other state or behavior of the system”
- TG: revisit in “composability criteria”

Most CXs are stateful [basis] (§1.3.1)

- CX custom instructions may be stateful
 - May access accumulators, registers, register files, ..
 - Conflict-free CX custom CSRs
- Uniform CX state context management
 - Supports 100s of CX state contexts, per CX, per system
 - Any mapping from harts to CX state contexts, 1:1, 1: n , n :1
 - Uniform OS access control, virtualization, context switching of any CX state context